

- KERN Language Development Plan
 - Project Overview
 - Core Objectives
 - Development Phases
 - Phase 1: Foundation (Grammar, AST, Validation)
 - 1.1 Lexer and Parser Implementation
 - 1.2 Abstract Syntax Tree (AST) Definition
 - 1.3 Static Validation
 - 1.4 Testing Framework
 - Phase 2: Rule Engine and Graph Builder (4-6 weeks)
 - 2.1 Execution Graph Data Model
 - 2.2 Rule Engine Implementation
 - 2.3 Flow Execution Engine
 - 2.4 Testing
 - Phase 3: Virtual Machine and Bytecode (6-8 weeks)
 - 3.1 Bytecode Instruction Set Implementation
 - 3.2 Virtual Machine Core
 - 3.3 Bytecode Compiler
 - 3.4 Performance and Security
 - 3.5 Testing
 - Phase 4: Tooling and Ecosystem (4-6 weeks)
 - 4.1 Development Tools
 - 4.2 Integration Tools
 - 4.3 Documentation and Examples
 - 4.4 Testing
 - Technical Architecture
 - Core Components
 - Data Flow
 - Performance Targets
 - Implementation Technologies
 - Risk Mitigation
 - Success Metrics

KERN Language Development Plan

Project Overview

KERN (Knowledge Execution & Reasoning Notation) is a deterministic intelligence execution language designed to encode business logic, rules, and workflows compactly. It's a logic-centric, graph-native execution language optimized for machine reasoning rather than human comfort.

Core Objectives

- Create a deterministic, rule-based execution language
- Build a graph-based execution model
- Ensure PSI.brain can easily parse, analyze, refactor, and generate KERN
- Achieve extreme performance with minimal storage requirements
- Maintain auditability and security

Development Phases

Phase 1: Foundation (Grammar, AST, Validation)

Duration: 4-6 weeks

1.1 Lexer and Parser Implementation

- Implement lexer based on EBNF grammar specification
- Create token definitions for all lexical elements
- Build recursive descent parser
- Handle all grammar productions (entities, rules, flows, constraints)
- Add comprehensive error reporting for syntax errors

1.2 Abstract Syntax Tree (AST) Definition

- Define AST node structures for all language constructs
- Create node types for entities, rules, flows, constraints
- Implement AST visitor pattern for traversal

- Add position tracking for debugging
- Design serialization/deserialization for PSI integration

1.3 Static Validation

- Implement type checking algorithms
- Create symbol table management
- Add scope resolution validation
- Implement dependency analysis
- Design rule conflict detection
- Create bytecode compatibility validation

1.4 Testing Framework

- Set up unit testing infrastructure
- Create test cases for all grammar productions
- Implement parser error recovery tests
- Add validation edge case tests

Phase 2: Rule Engine and Graph Builder (4-6 weeks)

2.1 Execution Graph Data Model

- Define graph node structures for operations, rules, conditions
- Implement edge definitions for data and control flow
- Create graph builder from AST
- Design graph optimization algorithms
- Implement cycle detection algorithms

2.2 Rule Engine Implementation

- Create rule matching algorithms
- Implement pattern matching engine
- Design rule priority system
- Build conflict resolution mechanisms
- Implement rule execution scheduling
- Add recursion prevention with explicit limits

2.3 Flow Execution Engine

- Implement flow pipeline execution
- Create demand-driven evaluation
- Build lazy evaluation strategies
- Design context passing mechanisms
- Implement control flow operations (if/then/else, loop, break, halt)

2.4 Testing

- Create rule evaluation test cases
- Test graph building algorithms
- Validate execution order
- Test rule conflict scenarios

Phase 3: Virtual Machine and Bytecode (6-8 weeks)

3.1 Bytecode Instruction Set Implementation

- Implement fixed-width instruction encoding (8 bytes per instruction)
- Create instruction decoder
- Build all control flow instructions (NOP, JMP, JMP_IF, HALT)
- Implement data and symbol instructions (LOAD_SYM, LOAD_NUM, MOVE, COMPARE)
- Add graph operations instructions
- Create rule execution instructions
- Implement context and state instructions
- Build error handling instructions
- Add external interface instructions

3.2 Virtual Machine Core

- Design register-based execution model (R0-R15, CTX, ERR, PC, FLAG)
- Implement instruction fetch-execute cycle
- Create memory management system
- Build context management
- Implement error handling as data

- Add introspection hooks for PSI
- Create step-by-step execution capability

3.3 Bytecode Compiler

- Build AST to bytecode translator
- Implement register allocation
- Create bytecode optimization passes
- Add bytecode serialization
- Implement bytecode verification

3.4 Performance and Security

- Implement memory limits
- Add execution step limits
- Create sandboxed execution environment
- Implement security validation
- Add performance monitoring

3.5 Testing

- Create bytecode execution tests
- Test VM instruction execution
- Validate performance targets (<10ms startup)
- Test security sandboxing

Phase 4: Tooling and Ecosystem (4-6 weeks)

4.1 Development Tools

- Create command-line compiler
- Build debugger with step execution
- Implement profiler
- Create bytecode inspector
- Build graph visualizer
- Add syntax highlighting support

4.2 Integration Tools

- Create external function adapters
- Implement serialization/deserialization tools
- Build PSI integration APIs
- Create import/export utilities

4.3 Documentation and Examples

- Write comprehensive language reference
- Create tutorial examples
- Build API documentation
- Create best practices guide

4.4 Testing

- End-to-end integration tests
- Performance benchmarking
- Tooling functionality tests
- User acceptance testing

Technical Architecture

Core Components

1. **Lexer/Parser:** Converts source code to AST
2. **Validator:** Ensures semantic correctness
3. **Graph Builder:** Creates execution graphs from AST
4. **Rule Engine:** Executes rule-based logic
5. **Bytecode Compiler:** Translates graphs to bytecode
6. **Virtual Machine:** Executes bytecode
7. **Tooling:** Development and debugging utilities

Data Flow

Source Code → Lexer → Parser → AST → Validator → Graph Builder → Bytecode Compiler → VM → Execution

Performance Targets

- Startup time: < 10ms
- Execution: Near-native performance
- Memory: Bounded and predictable
- Compilation: Fast and deterministic

Implementation Technologies

- **Language:** Rust/C++ for performance-critical components
- **Build System:** Cargo/CMake
- **Testing:** Built-in testing framework
- **Serialization:** Custom binary format for bytecode
- **Memory Management:** Manual or smart pointers

Risk Mitigation

- Follow KERN design principles strictly to avoid feature creep
- Implement comprehensive testing at each phase
- Maintain backward compatibility from early versions
- Regular PSI integration validation
- Performance monitoring throughout development

Success Metrics

- Successful compilation of all KERN grammar constructs
- Performance targets met
- PSI integration working seamlessly
- Security sandboxing effective
- Tooling providing good developer experience
- Deterministic execution guaranteed